



UNIVERSITÀ DI PISA



Scuola Superiore
Sant'Anna

di Studi Universitari e di Perfezionamento

A Framework for static allocation of parallel OpenMP code on multi-core platforms

Giacomo Dabisias, Filippo Brizzi

Università degli studi di Pisa,
Scuola Superiore Sant'Anna
Pisa, Italy

February 28, 2014



Context and motivations

Real-time systems are moving towards multicore architectures. The majority of multithread/core libraries target high performance systems.

- ▶ Real-time applications need strict timing guarantees and predictability.

Vs

- ▶ High performance systems try to achieve a lower computation time in a best effort manner.

There is no actual automatic tool which has the advantages of HPC with timing constraints.

Objectives

The developed framework has the following characteristics.

- ▶ An easy API to specify the concurrency between real- time tasks and scheduling parameters.
- ▶ A way to visualize task concurrency and code structure as graphs.
- ▶ A scheduling algorithm which supports multicore architectures, adapting to the specific platform.
- ▶ A run time support for the program execution which guarantees the scheduling order of tasks and their timing constrains.

Design Choice: OpenMP and Clang

OpenMP

- ▶ Minimal code overhead.
- ▶ Well spread standard.
- ▶ Opensource and supported by several vendors like Intel and IBM.

Clang

- ▶ Provides code analysis and source to source translation capabilities.
- ▶ Modularity and great efficiency.
- ▶ Opensource and supported by several vendors like Google and Apple.

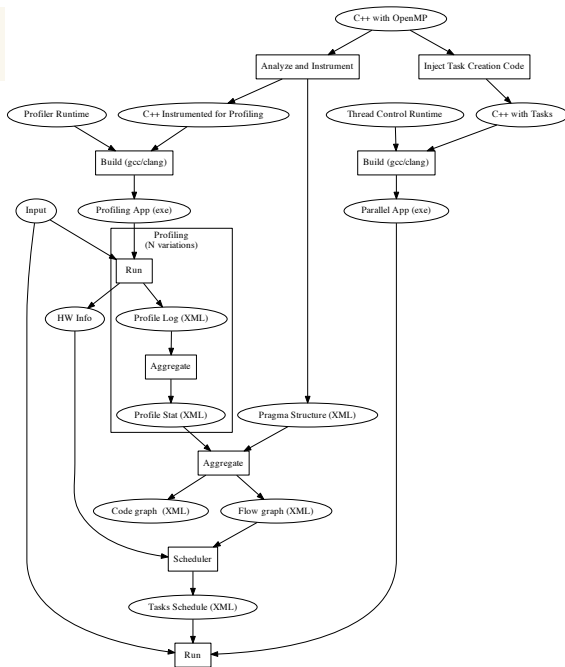
In July 2013 Intel released a patched version of Clang which fully supports the OpenMP 3.3 standard.



General Design

The framework takes as input a C++ code annotated with OpenMP.

- ▶ The pragmas are extracted with all relevant informations using Clang and saved as XML .
- ▶ The input code is rewritten to perform profiling .
- ▶ The scheduler tool uses these informations to create a possible schedule.
- ▶ The input code is rewritten to allow execution according to the generated schedule.
- ▶ The code is then executed with a custom run-time support.





Simple Example

```

1 void work(int bar){
2     #pragma omp parallel for
3     for (int i = 0; i < bar; ++i)
4     {
5         //do stuff
6     }
7 };
8 int main(int argc, char* argv[]) {
9     int bar;
10    #pragma omp parallel private(bar)
11    {
12        #pragma omp sections
13        {
14            #pragma omp section
15            {
16                //do stuff (bar)
17                work(bar);
18            }
19            #pragma omp section
20            {
21                //do stuff (bar)
22                work(bar);
23            }
24        }
25    }
26 }

```



OpenMP

Multiple threads of execution perform tasks defined by directives.

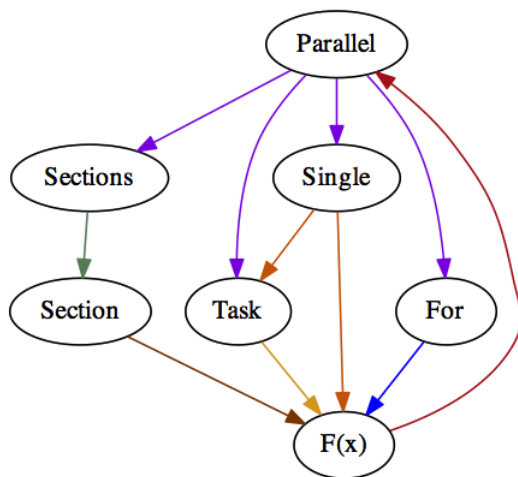
- ▶ Each directive applies to a block of C++ code embedded in a scope.
- ▶ Allows nested parallelism through nested directives.
- ▶ Clauses allow variables management.

```
#pragma omp directive-name [clause[ [,] clause]...] new-line
```

Chosen subset for the framework:

- ▶ Control directives : parallel, sections, single.
- ▶ Working directives : task, section, for.

OpenMP



Clang

Clang and OpenMP:

The strength of Clang lies in its implementation of the Abstract Syntax Tree (AST).

- ▶ Closely resembles both the written C++ code and the C++ standard.
- ▶ Clang's AST nodes are modeled on a class hierarchy that does not have a common ancestor.
- ▶ Hundreds of classes for a total of more than one hundred thousand lines of code.

Clang - AST

To traverse the AST, Clang provides the RecursiveASTVisitor class.

- ▶ Very powerful and easy to learn interface
- ▶ Possibility to create a custom visitor that triggers only on specific nodes.

Clang supports the insertion of custom code through the Rewriter class.

- ▶ Allows insertion, deletion and replacement of code.
- ▶ Operations are performed during the AST visit.
- ▶ A new source file with all the modifications is generated at the end of the visit.

Clang - AST

```

1 class A {
2 public:
3     int x;
4     void set_x(int val) {
5         x = val * 2;
6     }
7     int get_x() {
8         return x;
9     }
10 };
11 int main() {
12     A a;
13     int val = 5;
14     a.set_x(val);
15 }

```

TranslationUnitDecl

```

| - CXXRecordDecl <clang_ast_test.cpp:2:1, line:13:1>
|   class A
| | - CXXRecordDecl <line:2:1, col:7> class A
| | | - AccessSpecDecl <line:3:1, col:7> public
| | | - FieldDecl <line:4:2, col:6> x 'int'
| | | - CXXMethodDecl <line:5:2, line:7:2> set_x 'void_(
| | |   int)'
| | | | - ParmVarDecl <line:5:13, col:17> val 'int'
| | | | - CompoundStmt <col:22, line:7:2>
| | | |   - BinaryOperator <line:6:3, col:13> 'int' lvalue
| | | |     '='
| | | |       - MemberExpr <col:3> 'int' lvalue ->x
| | | |         - CXXThisExpr <col:3> 'class A*' this
| | | |         - BinaryOperator <col:7, col:13> 'int' '*'
| | | |         - ImplicitCastExpr <col:7> 'int' <
| | | |           LValueToRValue>
| | | |         - DeclRefExpr <col:7> 'int' lvalue ParmVar
| | | |           'val' 'int'
| | |         - IntegerLiteral <col:13> 'int' 2
| |
| ...

```

Instrumentation for Profile

- ▶ Creation of a custom profiler to time pragma code blocks and functions. No existing profiling tool allows this operation.
 - ▶ Code is instrumented with calls to a custom run-time support.
 - ▶ Execution time, children execution time, caller identifier, for loop counter.
 - ▶ Output is saved in an XML file.

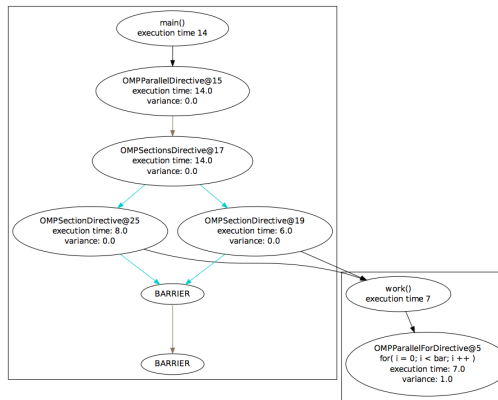
```
1 ...  
2 //#pragma omp parallel for  
3 if( ProfileTracker profile_tracker = ProfileTrackParams(3, 5, bar - 0))  
4 for (int i = 0; i < bar; ++i)  
5 {  
6     //do stuff  
7 }  
8 ...  
9 //#pragma omp section  
10 if( ProfileTracker profile_tracker = ProfileTrackParams(12, 25))  
11 {  
12     //do stuff (bar)  
13     work(bar);  
14 }  
15 ...
```

Graphs

- ▶ Pragas are extracted from the source file and saved in an XML file.
- ▶ Profiler is executed N times and statistics are obtained.

Using these informations two pragma graphs are created.

- ▶ Flow Graphs: illustrates the parallel dependencies between tasks.
- ▶ Code Graph: illustrates the pragma nesting.



Scheduler

Intrumentation for profiling - Annotated example

Flow graph

Scheduler

Scheduler - Search tree

Scheduler - Constraints check

Scheduler - (Cetto & Chetto)

Final execution



Final execution - Intrumentation



Final execution - Run-time



Fianl execution - (thread pool)

Final execution - (multiple job queues)



Final execution - (synchronization)



General structure

General structure -(graph of the test code)

Results

Results - (some tables and graphs)

Results - (total completion time)

Results - (service time - boxplot)

Results - (Jitter)

Results - Comments



